

Predicting Recovery Days

Woojin Justin Joseph Ben

Monday 15th December, 2025



In this notebook, we will be conducting some basic EDA and feature engineering to eventually construct a Multiple Linear Regression Model.

1 Import the necessary modules

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from utilities.functions import *
```

2 Load the Data

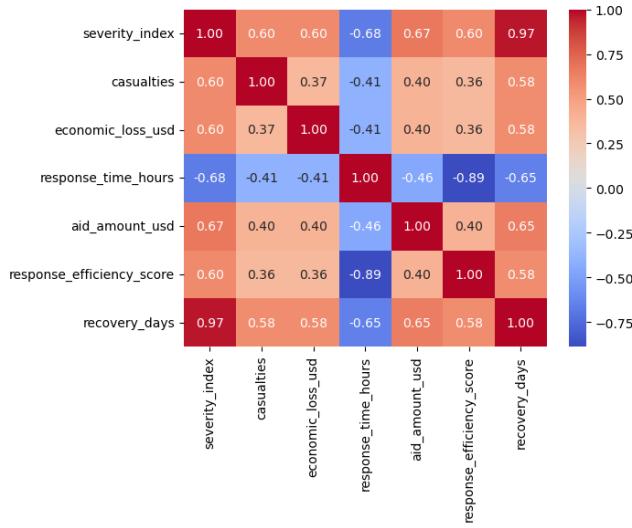
```
data = pd.read_csv('data/global_disaster_response_2018_2024 (1).csv')

data.head()
```

	date	country	disaster	severity_index	casualties	economic_loss_usd	response_time_hours	aid_amount_usd	response_efficiency_score	recovery_days	log_recovery_days	log_aid_amount_usd
0	2021-01-31	Brazil	Earthquake	0.99	111	7934365.62	2716033.21	67	-	-	30.613	22.557
1	2018-12-23	Brazil	Extreme Heat	1.65	100	8307648.09	265879.81	55	10.859	-	159.194	-
2	2020-08-10	India	Hurricane	5.5	22	7651362.94	49356.09	40.40	22	0.643	-	160.978
3	2022-09-15	Indonesia	Extreme Heat	1.65	94	13082518.31	237512.64	11.47	-	-	30.350	33.547
4	2022-09-28	United States	Wildfire	8.0	64	2655824.36	188910.88	1.42	-	-	19.170	17.137

2.1 Plotting a correlation heatmap to get an idea of the relationship between covariates and the response (recovery_days)

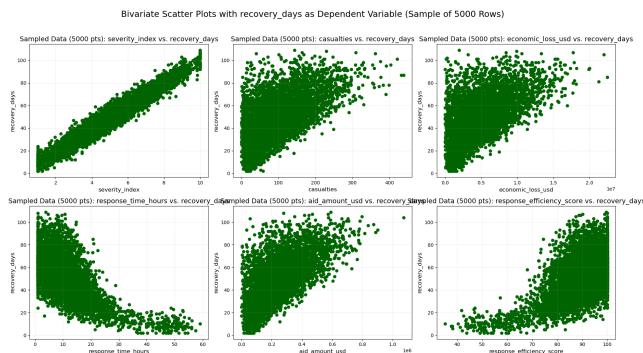
```
subset = data[['severity_index', 'casualties', 'economic_loss_usd', 'response_time_hours', 'aid_amount_usd', 'response_efficiency_score', 'recovery_days']]
corr_mat = subset.corr()
sns.heatmap(corr_mat, annot = True, fmt = ".2f", cmap = "coolwarm")
plt.savefig("figures/correlation_heatmap.png", bbox_inches='tight')
```



3 Check the relationships of different covariates with recovery_days to check linearity

Notice that we are plotting a random, representative subset of the data (5000 rows) as opposed to the entire dataset. Plotting the entire dataset would result in overplotting, and we'd have a harder time interpreting the plots.

```
plot_bivariate_scatter(
    data=data,
    y_column='recovery_days',
    x_columns=[
        'severity_index',
        'casualties',
        'economic_loss_usd',
        'response_time_hours',
        'aid_amount_usd',
        'response_efficiency_score'
    ],
    sample_size=5000
)
plt.savefig("figures/Bivariate_relationships.png", bbox_inches='tight')
```



4 Apply Log transformations on the some of the covariates to linearize its relationship with recovery_days

I'm also essentially inverting the meaning of response_efficiency_score as intuitively I thought it made more sense that as the score goes up, the recovery days would go down.

```
# Create a transformed copy
subset_tf = subset.copy()
```

```

# Apply transformations
subset_tf['severity_index_tf'] = subset_tf['severity_index'] # no transform

subset_tf['casualties_tf'] = np.log1p(subset_tf['casualties'])

subset_tf['economic_loss_usd_tf'] = np.log1p(subset_tf['economic_loss_usd'])

subset_tf['response_time_hours_tf'] = np.log1p(subset_tf['response_time_hours'])

subset_tf['aid_amount_usd_tf'] = np.log1p(subset_tf['aid_amount_usd'])

subset_tf['response_efficiency_score_tf'] = np.log1p(
    100 - subset_tf['response_efficiency_score']
)

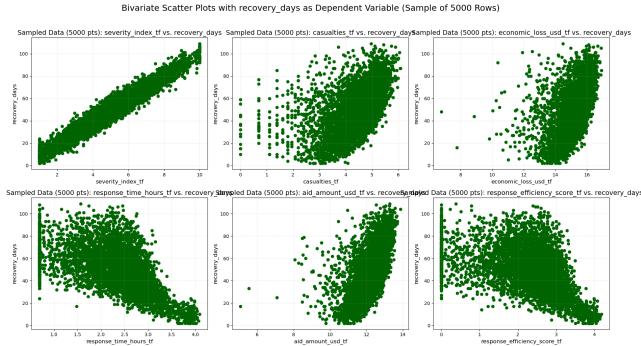
# Keep recovery_days unchanged
subset_tf['recovery_days'] = subset_tf['recovery_days']

y_column = 'recovery_days'

x_columns_tf = [
    'severity_index_tf',
    'casualties_tf',
    'economic_loss_usd_tf',
    'response_time_hours_tf',
    'aid_amount_usd_tf',
    'response_efficiency_score_tf'
]

plot_bivariate_scatter(
    data=subset_tf,
    y_column=y_column,
    x_columns= x_columns_tf,
    sample_size=5000
)
plt.savefig("figures/Bivariate_transformed_relationships.png", bbox_inches='tight')

```



We are observing some outliers towards the left of the plot, hence a good step would be to filter out the outliers depending on whether a data point is above or below 1.5 IQR away from the mean.

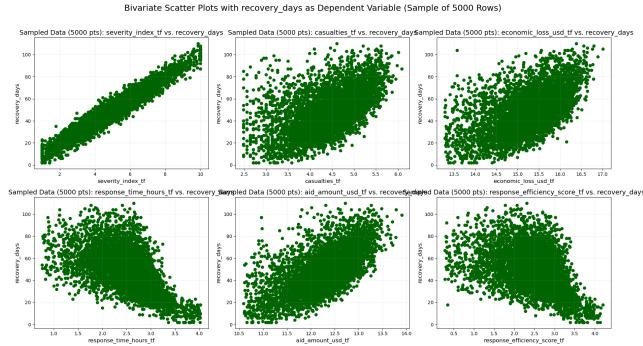
```
# Columns to check for outliers (transformed features only)
outlier_columns = [
    'severity_index_tf',
    'casualties_tf',
    'economic_loss_usd_tf',
    'response_time_hours_tf',
    'aid_amount_usd_tf',
    'response_efficiency_score_tf'
]

subset_tf_clean = remove_iqr_outliers(
    subset_tf,
    columns=outlier_columns,
    k=1.5
)

print(f"Rows before outlier removal: {len(subset_tf)}")
print(f"Rows after outlier removal: {len(subset_tf_clean)}")

Rows before outlier removal: 50000
Rows after outlier removal: 39545

plot_bivariate_scatter(
    data=subset_tf_clean,
    y_column=y_column,
    x_columns=x_columns_tf,
    sample_size=5000
)
plt.savefig("figures/Bivariate_relationships_wo_outliers.png", bbox_inches='tight')
```



5 Perform the train test split on the data, and fit the model on the transformed training design matrix.

```

X = data[
    [
        'severity_index',
        'casualties',
        'economic_loss_usd',
        'response_time_hours',
        'aid_amount_usd',
        'response_efficiency_score'
    ]
]

# Response variable
y = data['recovery_days']

features = [
    'severity_index',
    'casualties',
    'economic_loss_usd',
    'response_time_hours',
    'aid_amount_usd',
    'response_efficiency_score'
]

# Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

```

```

# Combine X and y for convenience
train_df = X_train.copy()
train_df['recovery_days'] = y_train

test_df = X_test.copy()
test_df['recovery_days'] = y_test

# Transform features
train_tf = transform_features(train_df)
test_tf = transform_features(test_df)

# Split back
X_train_tf = train_tf[features]
y_train_tf = train_tf['recovery_days']

X_test_tf = test_tf[features]
y_test_tf = test_tf['recovery_days']

# Fit
lin_reg = LinearRegression()
lin_reg.fit(X_train_tf, y_train_tf)

LinearRegression()

```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

[x]LinearRegression
[?Documentation for LinearRegression](#)iFitted

6 Predict using the test covariates and compare to the actual test data

```

y_pred = lin_reg.predict(X_test_tf)

# Metrics
print(f"R² = {r2_score(y_test_tf, y_pred):.4f}")
print(f"MSE = {mean_squared_error(y_test_tf, y_pred):.4f}")
print(f"MAE = {mean_absolute_error(y_test_tf, y_pred):.4f}")

R² = 0.9328
MSE = 25.0572
MAE = 3.9771

```

Based on the results above, our model seems to be doing a pretty good in explaining the variance in our response according to our features!

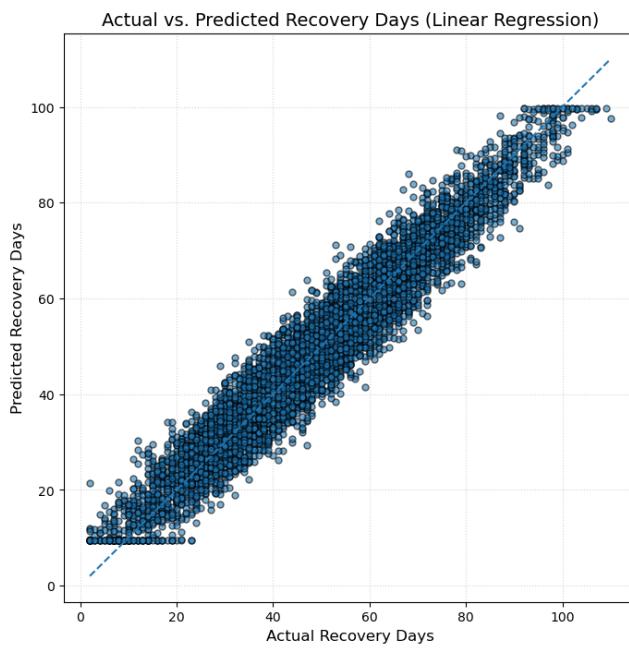
```
plt.figure(figsize=(8, 8))

plt.scatter(
    y_test_tf,
    y_pred,
    alpha=0.6,
    edgecolors='k',
    s=25
)

# 45 -degree reference line
min_val = min(y_test.min(), y_pred.min())
max_val = max(y_test.max(), y_pred.max())
plt.plot([min_val, max_val], [min_val, max_val], linestyle='--')

plt.xlabel("Actual Recovery Days", fontsize=12)
plt.ylabel("Predicted Recovery Days", fontsize=12)
plt.title("Actual vs. Predicted Recovery Days (Linear Regression)", fontsize=14)
plt.grid(True, linestyle=':', alpha=0.5)

plt.savefig("figures/actual_v_predicted.png", bbox_inches='tight')
```



Plotting the predicted vs actual recovery days, we can see our plot roughly following the line $y = x$, which is indicative of a model performing reasonably well in generalizing to the data.

```
residuals = y_test_tf - y_pred

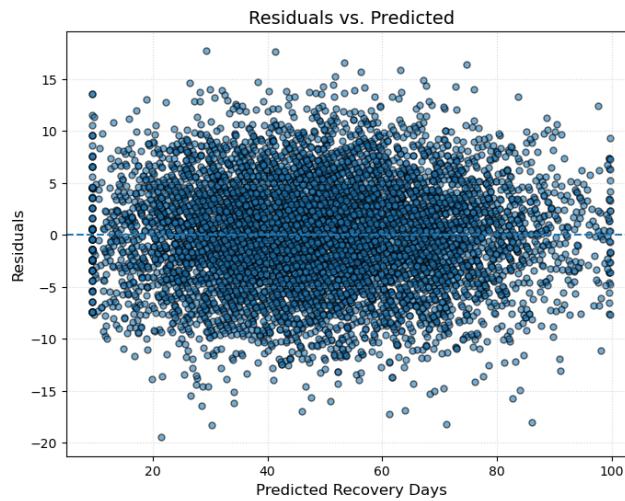
plt.figure(figsize=(8, 6))

plt.scatter(
    y_pred,
    residuals,
    alpha=0.6,
    edgecolors='k',
    s=25
)

plt.axhline(0, linestyle='--')

plt.xlabel("Predicted Recovery Days", fontsize=12)
plt.ylabel("Residuals", fontsize=12)
plt.title("Residuals vs. Predicted", fontsize=14)
plt.grid(True, linestyle=':', alpha=0.5)

plt.savefig("figures/residuals_v_predicted.png", bbox_inches='tight')
```



Plotting the predictions vs the residuals, we observe pretty much no heteroskedasticity, which is also indicative of our model generalizing to the data.